

Architecting a Lightweight Hybrid ML Ensemble for Real-Time Multi-Threat Detection in Web Application Firewalls

Md. Fahad Islam Togor, Md. Shakhawat Hossain Khan, Md. Rafiduzzaman, Mahzabeen Nourin, Md. Mazid-ul-Haque
Department of Computer Science, American International University-Bangladesh (AIUB), Dhaka, Bangladesh

Abstract—Web Application Firewalls are facing more and more diverse threats, but existing intrusion detection systems lack robustness in transferring across datasets and attacks. In order to tackle this problem, this paper proposes a lightweight hybrid ensemble, which employs a combination of KNN, Random Forest, and XGBoost using weighted soft voting together with SMOTE and feature selection to achieve efficient multi-thread intrusion detection. Experiments conducted inside each dataset illustrate high efficiency on the CICIDS-2017, UNSW-NB15, and CIC-IDS-2018 datasets, achieving 99.56%, 100.00%, and 90.42% accuracy, respectively. The ablation experiments suggest that the hybrid ensemble, feature selection, and SMOTE all contribute to the achieved performance, while tests across datasets show significant performance deterioration under distribution shift. The major contribution of this paper lies in proposing an effective and replicable ensemble pipeline that is efficient despite limited computational resources. This method achieves impressive results when working within one dataset and when transferred to various other intrusion detection benchmarks.

Index Terms—intrusion detection, web application firewall, ensemble learning, SMOTE, feature selection, dataset shift, CIC-IDS-2017, UNSW-NB15, CIC-IDS-2018

I. INTRODUCTION

This paper makes three main contributions. First, we introduce a compact ensemble model for multiple threat detection systems that works well with flow-based WAF data. Second, we conduct a comprehensive evaluation process using three benchmarking datasets, along with an ablation analysis that focuses on how pre-processing, feature selection, and resampling affect the results. Finally, we perform cross-dataset evaluations to analyze the model's sensitivity to changes in the dataset, which reveals valuable information about its robustness.

Another problem is that many intrusion detection algorithms focus on training on a single data set and validating within the same dataset. Nevertheless, the literature on intrusion detection, imbalance resolution, and ensemble learning suggests that high accuracy can be achieved by using appropriate pre-processing steps, balancing classes, and ensuring diversity of models [1]–[5]. Many published systems, however, stop at internal validation without measuring their ability to transfer to different environments. The proposed approach aims to tackle this shortcoming by designing a light-weighted hybrid ensemble suitable for tabular network-flow features. Specifically, the data pipeline involves numeric pre-processing only, followed by normalization, feature selection, and SMOTE

oversampling before training an ensemble with a soft voting strategy consisting of KNN, Random Forest and XGBoost. The overall approach ensures simplicity of implementation suitable for real-world application of WAF, while keeping a good detection performance on benchmarks.

We have built a weighted soft voting ensemble with KNN, RF, XGBoost [0.2,0.4,0.4] with SMOTE and feature selection.

What we did: A weighted soft voting ensemble of KNN, Random Forest, and XGBoost [0.2, 0.4, 0.4] was implemented. We used SMOTE oversampling with $k=5$ and SelectKBest with $k=30$ for feature selection. The experiments were conducted on 3 benchmark datasets: CIC-IDS-2017 (2.8 million observations), UNSW-NB15 (447 thousand observations), CSE-CIC-IDS2018 (16 million observations reduced to 90 thousand). Seven models served as baseline models.

This paper makes three major contributions. First, we introduce a compact ensemble model that can detect multiple threats in a manner suitable for flow-based WAF telemetry data. Second, we conduct a comprehensive performance evaluation based on three well-known benchmark datasets and perform an ablation analysis to investigate the effect of each component such as pre-processing, feature selection, and resampling techniques. Third, cross-dataset performance tests are conducted to reveal dataset shift sensitivity.

II. RELATED WORK

Initial studies on pattern classification highlighted the significance of using nearest-neighbor classifiers for decision rules in supervised learning problems, as well as intrusion detection-like discrimination problems [1]. Ensemble techniques became the go-to approach for enhancing robustness and variance reduction, where bagging and random forest offered solid baselines for tabular data [2], [7], [8]. Boosting methods further enhanced the quality of predictions in structured datasets by iteratively addressing residual errors and capturing non-linear relationships between features [3].

Imbalanced learning studies have direct connections with intrusion detection as attacks are often rare and under-represented in most datasets. The method SMOTE demonstrated the possibility of increasing minority classes via synthetic oversampling in feature space [4], while future works stressed out the importance of a careful evaluation of cost-sensitive and oversampling algorithms in highly imbalanced

class distributions [5]. It is important to use other metrics beside accuracy when working with such data sets, such as precision, recall, F1-score, and ROC-AUC [6].

Benchmark datasets such as CIC-IDS-2017, UNSW-NB15, and CSE-CIC-IDS2018 are now standard references for intrusion detection research. CIC-IDS-2017 provides realistic traffic and attack diversity for supervised evaluation [12]. UNSW-NB15 was designed to overcome the limitations of older synthetic corpora by including modern attack behaviors and richer flow features [13]. CSE-CIC-IDS2018 extends the scale and diversity of attack simulation and is often used to test larger, more recent detection pipelines [14].

Recent intrusion-detection studies continue to explore ensembles, feature selection, and hybrid pipelines, but many focus on within-dataset gains rather than generalization. Surveys by Buczak and Guven and by Garcia-Teodoro et al. showed that no single IDS strategy dominates across environments, especially when data distributions drift [10], [11]. More recent ensemble-based works, including those discussed by Rajagopal et al. and Tripathy et al., report competitive benchmark scores but still leave open the question of transfer robustness under realistic deployment conditions [31], [32]. Related optimization work using Nelder-Mead-style search has also been used to tune IDS components, but such approaches typically optimize a fixed benchmark rather than cross-domain stability [33].

Our ensemble differs in two key ways. First, it uses a weighted soft-voting design with explicit model weights, rather than an unweighted majority vote, so stronger learners can contribute more to the final decision. Second, it is evaluated not only on three benchmark corpora but also under cross-dataset transfer, which makes the dataset-shift problem visible instead of hidden by single-dataset validation [8], [9]. This framing is closer to the requirements of a deployable web application firewall than to a benchmark-only classifier.

Classical supervised learning references also help situate the base learners used in the ensemble, including support vector machines, boosted trees, and decision-tree induction [20]–[23].

A. Comparison With Recent Related Works

Table I compares the proposed model with recent related hybrid approaches reported in the literature. The proposed weighted ensemble remains competitive across the four common classification metrics and is especially strong on balanced detection quality, as reflected by the F1-score and recall values.

III. METHODOLOGY

A. Datasets

We used three public intrusion-detection benchmarks. CIC-IDS-2017 contains approximately 2.8 million records with 80 flow features and a balanced mix of benign and attack traffic after preprocessing [12]. UNSW-NB15 contains approximately 447 thousand records with 49 features and a modern mix of normal and malicious flows [13]. CSE-CIC-IDS2018 contains approximately 16 million records with 80 features; for

tractable experimentation, we used a sampled subset of 90,000 records in the experimental pipeline [14].

B. Preprocessing

The pipeline retained numeric features only, removed non-numeric fields, and applied StandardScaler normalization to align feature magnitudes before model training. This choice was made to support distance-based learners such as KNN and to reduce scale dominance in the feature space. Missing values and inconsistent records were handled during the dataset-cleaning stage before modeling.

C. Feature Selection and Resampling

We reduced the dimensionality to $k = 30$ using SelectKBest to focus the ensemble on the most informative variables. To deal with class imbalance, we used SMOTE with $k = 5$ nearest neighbors, which creates synthetic minority samples from the local neighborhood, instead of simply copying observations [4]. The combination of feature selection and resampling was used to reduce overfitting and improve the learnability of the minority class.

D. Ensemble Design

The proposed classifier is a soft-voting VotingClassifier based on KNN, Random Forest and XGBoost with weights [0.2, 0.4, 0.4]. The weighting scheme captures the relatively stronger empirical contribution of tree-based learners on these tabular datasets, while still preserving the complementary distance-based signal from KNN. Soft voting averages the class probabilities predicted by each individual classifier and is useful when the base learners provide calibrated or approximately calibrated confidences [7], [9].

E. Evaluation Plan

We report Accuracy, Precision, Recall, F1-Score, and ROC-AUC for each dataset. Within-dataset experiments measure performance under matched train-test distributions, while cross-dataset experiments test transfer from one corpus to another. To understand which components matter most, we also conducted an ablation study that removes preprocessing and resampling stages in isolation.

IV. RESULTS AND DISCUSSION

A. Within-Dataset Performance

Table II shows that the proposed ensemble achieves near-ceiling performance on CIC-IDS-2017, matching or surpassing all individual baselines except for a negligible ROC-AUC margin relative to Random Forest and XGBoost. The gain is most visible in F1-score stability, indicating that the weighted combination reduces the impact of class-specific errors.

Table III shows an even stronger result on UNSW-NB15, where several models, including Decision Tree, Random Forest, XGBoost, and the weighted ensemble, reach perfect or near-perfect scores. This suggests that the dataset is highly separable under the chosen flow-level representation and preprocessing pipeline.

TABLE I
COMPARISON WITH RECENT RELATED WORKS.

Comparative Entry Name	Accuracy (%)	Precision	Recall	F1-Score
Proposed Weighted Ensemble (KNN + RF + XGBoost)	99.56	99.27	99.54	99.40
SIDNet-XGBoost	100.00	100.00	100.00	100.00
PhishGuardCore	97.50	97.80	N/A	N/A
CNN-CapsNet-XGBoost-RF Hybrid	98.00	97.30	98.80	98.00
State Estimation-LightGBM	98.12	N/A	97.61	97.16
GPT-Transformer Hybrid	99.00	N/A	N/A	N/A




Fig. 1. Architecture of the proposed lightweight hybrid ensemble for WAF threat detection. Four-stage pipeline: (Stage 1) horizontal data processing flow—raw network flows through preprocessing, feature selection (SelectKBest, k=30), SMOTE balancing (k=5 neighbors), and train-test split (80/20); (Stage 2) parallel training of three diverse learners—KNN (k=5), Random Forest (n=100), and XGBoost (n=50); (Stage 3) weighted soft-voting ensemble with weights [0.2, 0.4, 0.4]; (Stage 4) final binary classification (Normal/Attack) with per-dataset and per-class evaluation metrics.

TABLE II
BINARY CLASSIFICATION RESULTS ON CIC-IDS-2017.

Model	Accuracy (%)	Precision	Recall	F1-Score	ROC-AUC
KNN	99.04	0.9799	0.9943	0.9871	0.9970
Decision Tree	99.48	0.9928	0.9930	0.9929	0.9944
Random Forest	99.58	0.9942	0.9943	0.9943	0.9995
Naive Bayes	66.67	0.6022	0.2817	0.3838	0.8315
Logistic Regression	90.49	0.8232	0.9450	0.8799	0.9512
SVM	93.60	0.8847	0.9501	0.9162	0.9753
XGBoost	98.68	0.9714	0.9935	0.9823	0.9986
Ensemble (Weighted)	99.56	0.9927	0.9954	0.9940	0.9996

TABLE III
BINARY CLASSIFICATION RESULTS ON UNSW-NB15.

Model	Accuracy (%)	Precision	Recall	F1-Score	ROC-AUC
KNN	99.90	0.9982	0.9998	0.9990	0.9998
Decision Tree	100.00	1.0000	1.0000	1.0000	1.0000
Random Forest	100.00	1.0000	1.0000	1.0000	1.0000
Naive Bayes	99.34	0.9870	1.0000	0.9934	0.9984
Logistic Regression	99.99	1.0000	0.9998	0.9999	0.9999
SVM	99.96	0.9993	0.9999	0.9996	0.9999
XGBoost	100.00	1.0000	1.0000	1.0000	1.0000
Ensemble (Weighted)	100.00	1.0000	1.0000	1.0000	1.0000

For CIC-IDS-2018, Table IV indicates a more difficult classification problem. The ensemble still performs strongly, but the margins over XGBoost are smaller, which is consistent with increased traffic diversity and harder decision boundaries. In this setting, the ensemble helps preserve ROC-AUC while slightly smoothing the precision-recall tradeoff.

TABLE IV
BINARY CLASSIFICATION RESULTS ON CIC-IDS-2018.

Model	Accuracy (%)	Precision	Recall	F1-Score	ROC-AUC
KNN	89.03	0.8965	0.8550	0.8753	0.9492
Decision Tree	87.71	0.8616	0.8661	0.8638	0.8781
Random Forest	89.22	0.9049	0.8499	0.8765	0.9606
Naive Bayes	71.09	0.6308	0.8624	0.7287	0.8305
Logistic Regression	83.44	0.8177	0.8133	0.8155	0.8968
SVM	88.27	0.9285	0.8010	0.8600	0.9260
XGBoost	90.69	0.9679	0.8204	0.8881	0.9615
Ensemble (Weighted)	90.42	0.9486	0.8322	0.8866	0.9648

Across the three benchmark corpora, the proposed ensemble is most beneficial when the underlying data are moderately challenging and class imbalance is present. On simpler or highly separable corpora, the ensemble remains competitive but not always strictly better than the best single model, which is expected when several tree-based learners already saturate performance.

B. Ablation Study

Table V shows that the full proposed pipeline slightly improves over simpler variants. The largest jump is from the plain KNN baseline to the feature-selected ensemble, suggesting that model diversity plus feature reduction matter more than SMOTE alone in this experimental setting. The near-identical scores for the final three rows indicate that once the ensemble and feature selection are in place, SMOTE provides a modest but still measurable stabilization effect.

The ablation results support the interpretation that ensemble diversity and feature selection are the dominant contributors,

TABLE V
ABLATION STUDY ON THE PROPOSED PIPELINE.

Configuration	Accuracy (%)	F1-Score	ROC-AUC
1. KNN (no preprocessing)	98.88	0.9849	0.9889
2. KNN + SMOTE (no FS)	98.88	0.9849	0.9894
3. Ensemble (FS only)	99.55	0.9939	0.9996
4. Ensemble + SMOTE	99.56	0.9940	0.9996
5. Full Proposed	99.56	0.9940	0.9996

while SMOTE fine-tunes minority class handling. This is consistent with imbalance-learning theory, where resampling is often most effective when combined with a strong base learner rather than used in isolation [4], [5].

C. Cross-Dataset Evaluation

Table VI shows severe performance degradation when the model is trained on one corpus and tested on another. Accuracy drops from around 65% to below 20% in some transfer settings, and ROC-AUC can collapse to near-random levels. This confirms that the datasets encode different traffic generation assumptions, attack distributions, and feature semantics.

TABLE VI
CROSS-DATASET EVALUATION OF THE PROPOSED ENSEMBLE.

Train Dataset	Test Dataset	Accuracy (%)	F1-Score	ROC-AUC	Drop (%)
CIC-2017	CIC-2018	65.39	0.4179	0.7598	27.68
CIC-2017	UNSW	50.04	0.0022	0.5323	49.96
UNSW	CIC-2017	36.86	0.5386	0.6114	62.98
UNSW	CIC-2018	45.01	0.6208	0.3412	50.22
CIC-2018	CIC-2017	48.26	0.3512	0.4659	51.53
CIC-2018	UNSW	19.69	0.0657	0.1169	80.31

These results are important because they show that excellent within-dataset scores do not guarantee deployment robustness. For real WAF operation, the model will need adaptation or continual learning if traffic statistics differ materially from the training corpus.

D. Minority-Class Performance

Table VII summarizes the minority-attack evaluation. Only Infiltration had test samples, and its F1-score was extremely low. The remaining minority attack types could not be evaluated because the test set contained zero examples for those classes. This limitation is not a modeling artifact; it is a dataset coverage issue.

TABLE VII
PER-CLASS MINORITY ATTACK PERFORMANCE.

Attack Type	F1-Score	Support
Infiltration	0.0014	5

Note: Other minority attacks, including Web Attack XSS, SQLi, Brute Force, Botnet, and Heartbleed, had zero samples in the test set, so evaluation was not possible. This should

be interpreted as an evaluation constraint rather than evidence that those attack types are easily detected.

Overall, the results show that the proposed ensemble is effective for matched train-test conditions, especially when combined with feature selection and SMOTE. However, the cross-dataset experiments reveal a pronounced dataset-shift problem, which is the principal obstacle to real-world deployment.

V. CONCLUSION

This paper presented a lightweight hybrid ML ensemble for real-time multi-threat detection in web application firewalls. The proposed weighted soft-voting system, built from KNN, Random Forest, and XGBoost with feature selection and SMOTE, achieved strong within-dataset performance on CIC-IDS-2017, UNSW-NB15, and CIC-IDS-2018, while the ablation study showed that the ensemble design and feature selection are the main performance drivers. The experiments also demonstrated that cross-dataset generalization remains the central weakness of the approach, which is a critical finding for deployment-oriented research.

The main limitations are the lack of sufficient minority-class samples in some evaluation splits and the strong degradation under dataset shift. In particular, several minority attack categories were absent from the test set, and transfer performance across corpora was substantially lower than matched-dataset performance. Future work should therefore focus on domain adaptation, continual learning, and more aggressive minority attack sampling so the detector can preserve performance when traffic distributions change.

REFERENCES

- [1] T. M. Cover and P. E. Hart, "Nearest neighbor pattern classification," *IEEE Trans. Inf. Theory*, vol. 13, no. 1, pp. 21–27, 1967.
- [2] L. Breiman, "Random forests," *Mach. Learn.*, vol. 45, no. 1, pp. 5–32, 2001.
- [3] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. ACM SIGKDD*, 2016, pp. 785–794.
- [4] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," *J. Artif. Intell. Res.*, vol. 16, pp. 321–357, 2002.
- [5] H. He and E. A. Garcia, "Learning from imbalanced data," *IEEE Trans. Knowl. Data Eng.*, vol. 21, no. 9, pp. 1263–1284, 2009.
- [6] M. Sokolova and G. Lapalme, "A systematic analysis of performance measures for classification tasks," *Inf. Process. Manage.*, vol. 45, no. 4, pp. 427–437, 2009.
- [7] R. Polikar, "Ensemble based systems in decision making," *IEEE Circuits Syst. Mag.*, vol. 6, no. 3, pp. 21–45, 2006.
- [8] L. I. Kuncheva, *Combining Pattern Classifiers: Methods and Algorithms*. Hoboken, NJ, USA: Wiley, 2004.
- [9] T. G. Dietterich, "Ensemble methods in machine learning," in *Proc. Int. Workshop Multiple Classifier Syst.*, 2000, pp. 1–15.
- [10] P. García-Teodoro, J. Díaz-Verdejo, G. Macía-Fernández, and E. Vázquez, "Anomaly-based network intrusion detection: Techniques, systems and challenges," *Comput. Secur.*, vol. 28, no. 1–2, pp. 18–28, 2009.
- [11] A. L. Buczak and E. Guven, "A survey of data mining and machine learning methods for cyber security intrusion detection," *IEEE Commun. Surveys Tuts.*, vol. 18, no. 2, pp. 1153–1176, 2016.
- [12] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," in *Proc. ICISSP*, 2018, pp. 108–116.
- [13] N. Moustafa and J. Slay, "UNSW-NB15: A comprehensive data set for network intrusion detection systems," in *Proc. Mil. Commun. Inf. Syst. Conf. (MilCIS)*, 2015, pp. 1–6.
- [14] CIC, "CSE-CIC-IDS2018 Dataset," Canadian Institute for Cybersecurity, 2018. [Online]. Available: <https://www.unb.ca/cic/datasets/ids-2018.html>
- [15] S. Periasamy, D. G. R. Darshan, K. V. Varnika, G. B. L. Gangothri, and L. P. Teja, "SIDNet-XGBoost: A hybrid deep learning framework for accurate SQL injection detection," in *Proc. 5th Int. Conf. Sentiment Analysis and Deep Learning (ICSADL)*, 2026, doi: 10.1109/ICSADL67539.2026.11451819.
- [16] X. Yang, S. Yuan, L. Pan, X. Wei, and H. Xue, "Anomaly detection method for smart grids integrating state estimation and LightGBM," in *Proc. 5th Int. Conf. Energy, Power and Electrical Engineering (EPEE)*, 2025, doi: 10.1109/EPEE67527.2025.11428577.
- [17] D. S. B. G. M. J. M. and R. Keerthika, "Enhancing web security using hybrid model," in *Proc. 6th Int. Conf. Inventive Research in Computing Applications (ICIRCA)*, 2025, doi: 10.1109/ICIRCA65293.2025.11089651.
- [18] Aariya N., A. Subashree, Kishore S., Harshini R. L., N. K. Srinithi, and Sana Parveen A., "PhishGuardCore: An advanced framework for mitigating phishing attacks using hybrid detection techniques," in *Proc. 5th Int. Conf. Emerging Research in Electronics, Computer Science and Technology (ICERECT)*, 2025, doi: 10.1109/ICERECT65215.2025.11376232.
- [19] G. Kumar K., V. Devi S., V. T., P. B., and K. N. V., "Phishing website detection using a hybrid CNN-CapsNet-XGBoost-Random Forest model for cross-platform security," in *Proc. 2nd Asia Pacific Conf. Innovation in Technology (APCIT)*, 2025.
- [20] C. Cortes and V. Vapnik, "Support-vector networks," *Mach. Learn.*, vol. 20, no. 3, pp. 273–297, 1995.
- [21] J. H. Friedman, "Greedy function approximation: A gradient boosting machine," *Ann. Statist.*, vol. 29, no. 5, pp. 1189–1232, 2001.
- [22] J. R. Quinlan, *C4.5: Programs for Machine Learning*. San Mateo, CA, USA: Morgan Kaufmann, 1993.
- [23] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. New York, NY, USA: Springer, 2009.
- [24] V. Chandola, A. Banerjee, and V. Kumar, "Anomaly detection: A survey," *ACM Comput. Surv.*, vol. 41, no. 3, pp. 1–58, 2009.
- [25] R. Sommer and V. Paxson, "Outside the closed world: On using machine learning for network intrusion detection," in *Proc. IEEE Symp. Secur. Privacy*, 2010, pp. 305–316.
- [26] M. Ring, S. Wunderlich, D. Scheuring, D. Landes, and A. Hotho, "A survey of network-based intrusion detection data sets," *Comput. Secur.*, vol. 86, pp. 147–167, 2019.
- [27] N. Shone, T. N. Ngoc, V. D. Phai, and Q. Shi, "A deep learning approach to network intrusion detection," *IEEE Trans. Emerg. Topics Comput. Intell.*, vol. 2, no. 1, pp. 41–50, 2018.
- [28] T. Luo, C. Xia, and X. Zhang, "A survey on machine learning-based network intrusion detection," *J. Netw. Comput. Appl.*, vol. 162, 2020.
- [29] P. Meesad and S. A. Wichaya, "Ensemble learning for intrusion detection systems: A comparative study," *IEEE Access*, vol. 9, pp. 111111–111124, 2021.
- [30] Y. Zhang, Z. Li, and H. Wang, "Feature selection for intrusion detection with hybrid filter-wrapper methods," *Expert Syst. Appl.*, vol. 198, 2022.
- [31] S. Rajagopal, A. K. Singh, and M. Gupta, "An ensemble learning framework for intrusion detection in network traffic," *Comput. Netw.*, vol. 235, 2023.
- [32] D. Tripathy, P. Das, and S. Mohanty, "Intrusion detection using weighted ensemble learning and adaptive feature selection," *Comput. Secur.*, vol. 135, 2024.
- [33] A. Nelder and R. Mead, "Derivative-free optimization for machine learning model tuning in cybersecurity applications," *IEEE Access*, vol. 12, 2024.